

Homework 2 - Advanced Database Systems (ICS424)

Alfaifi, Ammar – 201855360

March 27, 2024

First, to setup the and install the DBMS, choosing PostgreSQL, on macOS, see

```
% brew install postgresql@16
==> Downloading https://formulae.brew.sh/api/formula.json
==> Downloading https://formulae.brew.sh/api/cask.json
Warning: postgresql@16 16.2.1 is already installed and up-to-date.
To reinstall 16.2.1, run:
  brew reinstall postgresql@16
~ took 4s
% psql -l
```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
postgres	ammar	UTF8	libc	C	C			
template0	ammar	UTF8	libc	C	C			=c/ammar +
template1	ammar	UTF8	libc	C	C			=c/ammar +

(3 rows)

Figure 1: Install PostgreSQL package and checking it works by executing `psql -l`, to list all existing databases.

1 Startup and Shutdown Options and Procedures

Starting PostgreSQL on Linux/Unix-based systems as in Figure 2.

```
% brew services stop postgresql@16
Stopping `postgresql@16`... (might take a while)
==> Successfully stopped `postgresql@16` (label: homebrew.mxcl.postgresql@16)

% brew services start postgresql@16
==> Successfully started `postgresql@16` (label: homebrew.mxcl.postgresql@16)
```

Figure 2: Start and stop PostgreSQL.

2 Logical and Physical Structures of a Database

The logical structure of the database is represented by the schema, which defines the tables, relationships, constraints, and other logical entities. The physical structure is how the data is stored on disk, which includes file organization, indexing, and storage parameters.

We can view the logical structure of the Horse database by running the following command:

```
1 \d
```

See Figure 3 for list of all tables of Horse DB, after I fixed all issues of compatibility with PostgreSQL-16.

We can also inspect the physical storage parameters using the following command:

```
1 SELECT relname, relkind, relpages
2 FROM pg_catalog.pg_class
3 WHERE relkind IN ('r', 'i', 't');
```

```
postgres=# \d
```

List of relations			
Schema	Name	Type	Owner
public	horse	table	ammar
public	owner	table	ammar
public	owns	table	ammar
public	race	table	ammar
public	raceresults	table	ammar
public	stable	table	ammar
public	track	table	ammar
public	trainer	table	ammar

(8 rows)

Figure 3: Listing all tables in the DB.

This query will display the names of relations (tables, indexes, and toast tables), their type, and the number of disk pages they occupy, see Figure 4.

relname	relkind	relpages
race_pkey	i	1
track	r	0
race	r	0
raceresults_pkey	i	2
raceresults	r	1
stable_pkey	i	1
pg_statistic	r	10
pg_type	r	15
Horse_pkey	i	1
stable	r	0
Horse	r	0
owner_pkey	i	1
owns_pkey	i	1
owner	r	0
owns	r	0
trainer_pkey	i	1
trainer	r	0
track_pkey	i	1

Figure 4: The names of relations (tables, indexes, and toast tables), their type, and the number of disk pages they occupy.

3 Creating a Materialized View with a Complex Join

A materialized view is a pre-computed result set that is stored on disk for faster query performance. Here's an example of creating a materialized view that joins multiple tables:

```
1 CREATE MATERIALIZED VIEW horse_owner_stable_view AS
2 SELECT h.horseName, o.lname, o.fname, s.stableName, s.location
3 FROM Horse h
4 JOIN Owns ow ON h.horseId = ow.horseId
5 JOIN Owner o ON ow.ownerId = o.ownerId
6 JOIN Stable s ON h.stableId = s.stableId;
```

This materialized view combines data from the Horse, Owns, Owner, and Stable tables to provide a consolidated view of horse names, owner names, and stable information. See Figure 5.

4 Developing and Demonstrating a Stored Procedure and a Trigger

To create a Stored Procedure we run this command:

```
1 CREATE OR REPLACE FUNCTION update_horse_age()
2 RETURNS TRIGGER AS $$
3 BEGIN
4     UPDATE Horse
5     SET age = age + 1
6     WHERE horseId = NEW.horseId;
7     RETURN NEW;
```

```

postgres=# CREATE MATERIALIZED VIEW horse_owner_stable_view AS
SELECT h.horseName, o.lname, o.fname, s.stableName, s.location
FROM Horse h
JOIN Owns ow ON h.horseId = ow.horseId
JOIN Owner o ON ow.ownerId = o.ownerId
JOIN Stable s ON h.stableId = s.stableId;
postgres=# select * from horse_owner_stable_view;
horse_owner_stable_view
information_schema. public.         raceresults
stable
postgres=# select * from horse
horse
horse_owner_stable_view
postgres=# select * from horse_owner_stable_view;
horse
horse_owner_stable_view

```

horse	owner	race	track	
horse_name	owner	race	track	
warrior	sulaiman	zobair farm	riyadh	
conquerer	mohammed	dubai stables	dubai	
dove of peace	mohammed	khalid	zobair farm	riyadh
ever faster	mohammed	khalid	zobair farm	riyadh
slow winner	saeed	ahmed	zayed farm	dubai
slow winner	faleh	mahmood	zayed farm	dubai
slow winner	sulaiman	zayed farm	dubai	dubai
windrunner	saeed	ahmed	zayed farm	dubai
windrunner	nasr	zayed farm	dubai	dubai
catapult	mohammed	naeem	dubai stables	dubai
flying force	saleh	fahd	sunny stables	jubail
laggard	mohammed	khalid	sunny stables	jubail
formula one	nazir	mohammed	zayed farm	dubai
formula one	mohammed	faisal	zayed farm	dubai
frisky frolic	fahd	abdul rahman	sunny stables	jubail
fantastic	ahmed	faisal	zayed farm	dubai
midnight	saud	mohammed	zayed farm	dubai
running wild	saeed	ali	zayed farm	dubai
fastoffmyfeet	saeed	ali	zobair farm	riyadh
slow poke	mohammed	naeem	zayed farm	dubai
slinger	fahd	abdul rahman	zayed farm	dubai
front runner	mohammed	sheikh	sunny stables	jubail
night	faleh	mahmood	zobair farm	riyadh
negative	abed	ahmed	zayed farm	dubai
lightening	abed	ahmed	dubai stables	dubai
lazy loser	faisal	khan	zobair farm	riyadh
leaping lizard	faleh	mahmood	zobair farm	riyadh
beautiful brown	yahya	mohammed	dubai stables	dubai
sick winner	mohammed	khalid	zayed farm	dubai
lazy loser	said	sheikh	zobair farm	riyadh
sublime	mashour	dubai stables	dubai	dubai

(31 rows)

Figure 5: Creating a view and listing all of its contents.

```

8 END;
9 $$ LANGUAGE plpgsql;
10
11 CREATE TRIGGER update_horse_age_trigger
12 AFTER INSERT ON RaceResults
13 FOR EACH ROW
14 EXECUTE FUNCTION update_horse_age();

```

This stored procedure `update_horse_age()` is triggered after inserting a new row into the `RaceResults` table. It updates the age of the horse by incrementing it by 1 year. Now, to create a Trigger, we run instead this commmadn:

```

1 CREATE OR REPLACE FUNCTION prevent_duplicate_owners()
2 RETURNS TRIGGER AS $$
3 DECLARE
4     owner_count INT;
5 BEGIN
6     SELECT COUNT(*) INTO owner_count
7     FROM Owns
8     WHERE ownerId = NEW.ownerId AND horseId = NEW.horseId;
9     IF owner_count > 0 THEN
10         RAISE EXCEPTION 'Owner already owns this horse';
11     END IF;
12     RETURN NEW;
13 END;
14 $$ LANGUAGE plpgsql;
15

```

```

postgres=# BEGIN;
UPDATE Horse SET age = 3 WHERE horseId = 'horse1';
BEGIN
UPDATE 1
postgres=# BEGIN;
UPDATE Horse SET age = 4 WHERE horseId = 'horse1';
WARNING:  there is already a transaction in progress
BEGIN
UPDATE 1

```

Figure 6: When we run two sessions we get a warning indicating there's an already running transaction in progress.

```

postgres=# SELECT xmin, xmax, cmin, cmax, ctid
FROM pg_catalog.pg_class
WHERE relname = 'Horse';
 xmin | xmax | cmin | cmax | ctid
-----+-----+-----+-----+-----
(0 rows)
postgres=#

```

Figure 7: Here we see an empty list of the the timestamp table.

```

16 CREATE TRIGGER prevent_duplicate_owners_trigger
17 BEFORE INSERT ON Owns
18 FOR EACH ROW
19 EXECUTE FUNCTION prevent_duplicate_owners();

```

This trigger `prevent_duplicate_owners()` is executed before inserting a new row into the `Owns` table. It checks if the owner already owns the horse being inserted, and if so, it raises an exception to prevent the insertion.

5 Demonstrating Locking and Timestamping

PostgreSQL provides various locking mechanisms to ensure data integrity and concurrency control. We can observe locking by running two separate sessions and performing conflicting operations. For Session1:

```

1 BEGIN;
2 UPDATE Horse SET age = 3 WHERE horseId = 'horse1';

```

and for Session2:

```

1 BEGIN;
2 UPDATE Horse SET age = 4 WHERE horseId = 'horse1';

```

In Session 2, the `UPDATE` statement will be blocked until Session 1 commits or rolls back its transaction, demonstrating row-level locking. See Figure 6 for how we can do two transactions.

PostgreSQL also supports timestamping to track when data was last modified. We can view the timestamp of a table by running:

```

1 SELECT xmin, xmax, cmin, cmax, ctid
2 FROM pg_catalog.pg_class
3 WHERE relname = 'Horse';

```

See Figure 7 for the list of existing timestamps.

```

postgres=# SELECT xmin, xmax, cmin, cmax, ctid
FROM pg_catalog.pg_class
WHERE relname = 'Horse';
  xmin | xmax | cmin | cmax | ctid
-----+-----+-----+-----+-----
(0 rows)

postgres=#

```

Figure 8: PostgreSQL lock table

```

postgres=# BEGIN;
UPDATE Horse SET age = 4 WHERE horseid = 'horse1';
BEGIN
UPDATE t
postgres=# SELECT pid, username, application_name, client_addr, query
FROM pg_stat_activity;
 pid | username | application_name | client_addr | query
-----+-----+-----+-----+-----
 7560 |          |                  |             |
 7561 | ammar    |                  |             |
38499 | ammar    | psql              |             | SELECT pid, username, application_name, client_addr, query+
      |          |                  |             | FROM pg_stat_activity;
 7557 |          |                  |             |
 7556 |          |                  |             |
 7559 |          |                  |             |
(6 rows)

postgres=# select pg_terminate_backend(38499);
FATAL: terminating connection due to administrator command
server closed the connection unexpectedly
        This probably means the server terminated abnormally
        before or while processing the request.
The connection to the server was lost. Attempting reset: Succeeded.
postgres=# select pg_terminate_backend(38499);
WARNING: PID 38499 is not a PostgreSQL backend process
pg_terminate_backend
-----
f
(1 row)

postgres=# BEGIN;
UPDATE Horse SET age = 4 WHERE horseid = 'horse1';
BEGIN
UPDATE t
postgres=#

```

Figure 9: We searched up for a running process's PID and then terminate it.

6 Diagnosing and Resolving Locking Conflicts

To diagnose locking conflicts, We can use the `pg_locks` view, which provides information about outstanding locks in the database. For example:

```

1 SELECT l.locktype, l.relation::regclass, l.virtualtransaction, l.pid
2 FROM pg_locks l
3 JOIN pg_stat_activity a ON l.pid = a.pid;

```

This query displays the lock type, relation (table) name, virtual transaction ID, and process ID of the locks, as seen in Figure 8.

To resolve locking conflicts, we can either wait for the blocking transaction to complete or terminate the blocking process using the `pg_terminate_backend()` function. See Figure 9 for how we can terminate processes from their PID.

7 Evaluating Backup Options

PostgreSQL offers several backup options, including:

- **SQL Dump:** This creates a plain-text file containing SQL statements to recreate the database objects and data. We can create a SQL dump using the `pg_dump` utility.
- **File System Level Backup:** This involves creating a consistent backup of the PostgreSQL data directory using file system-level tools like `tar` or disk-level backup utilities.
- **Continuous Archiving and Point-in-Time Recovery (PITR):** This technique involves archiving Write-Ahead Log (WAL) segments, which can be used to restore the database to a specific point in time.

```

postgres=# CREATE INDEX idx_horse_age ON Horse (age);
CREATE INDEX
postgres=# SELECT relname, relkind, indkey
FROM pg_catalog.pg_index
JOIN pg_catalog.pg_class ON pg_index.indexrelid = pg_class.oid
WHERE relname LIKE 'idx_%';
   relname   | relkind | indkey
-----+-----+-----
idx_horse_age | i       | 3
(1 row)

postgres=# DROP INDEX idx_horse_age;
DROP INDEX
postgres=# SELECT relname, relkind, indkey
FROM pg_catalog.pg_index
JOIN pg_catalog.pg_class ON pg_index.indexrelid = pg_class.oid
WHERE relname LIKE 'idx_%';
   relname   | relkind | indkey
-----+-----+-----
(0 rows)

```

Figure 10: We created an index on `age` column, view it, and drop at the end as demonstration.

We can evaluate and choose the appropriate backup option based on our requirements, such as backup size, restore time, and the need for point-in-time recovery.

8 Recovering a Database

To recover a database from a backup, follow these general steps:

1. Stop the PostgreSQL server.
2. Replace the data directory with the backup files (for file system level backup) or restore the SQL dump (for SQL dump backup).
3. Start the PostgreSQL server.
4. If using PITR, restore the database from the archived WAL segments to the desired point in time.

9 Creating and Managing Indexes

Indexes improve query performance by allowing faster data retrieval. To create an index on a table column, use the `CREATE INDEX` statement:

```
1 CREATE INDEX idx_horse_age ON Horse (age);
```

This creates an index named `idx_horse_age` on the `age` column of the `Horse` table.

We can view existing indexes using the following query:

```

1 SELECT relname, relkind, indkey
2 FROM pg_catalog.pg_index
3 JOIN pg_catalog.pg_class ON pg_index.indexrelid = pg_class.oid
4 WHERE relname LIKE 'idx_%';

```

To drop an index, use the `DROP INDEX` statement:

```
1 DROP INDEX idx_horse_age;
```

See Figure 10, how apply all above in the `age` column in `Horse` table.

```

postgres=#
postgres=# EXPLAIN ANALYZE
SELECT h.horseName, o.lname, o.fname
FROM Horse h
JOIN Owns ow ON h.horseId = ow.horseId
JOIN Owner o ON ow.ownerId = o.ownerId
WHERE h.age > 3;

QUERY PLAN

-----
Hash Join (cost=22.01..40.84 rows=225 width=144) (actual time=0.098..0.123 rows=10 loops=1)
  Hash Cond: ((ow.ownerid)::text = (o.ownerid)::text)
    -> Hash Join (cost=1.44..19.68 rows=225 width=96) (actual time=0.053..0.065 rows=10 loops=1)
      Hash Cond: ((ow.horseid)::text = (h.horseid)::text)
      -> Seq Scan on owns ow (cost=0.00..16.50 rows=650 width=96) (actual time=0.007..0.010 rows=31 loops=1)
      -> Hash (cost=1.32..1.32 rows=9 width=96) (actual time=0.029..0.029 rows=8 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 9kB
        -> Seq Scan on horse h (cost=0.00..1.32 rows=9 width=96) (actual time=0.014..0.017 rows=8 loops=1)
          Filter: (age > 3)
          Rows Removed by Filter: 18
    -> Hash (cost=14.70..14.70 rows=470 width=144) (actual time=0.018..0.018 rows=20 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 10kB
      -> Seq Scan on owner o (cost=0.00..14.70 rows=470 width=144) (actual time=0.007..0.010 rows=20 loops=1)
Planning Time: 0.979 ms
Execution Time: 0.172 ms
(15 rows)

```

Figure 11: An execution plan example.

10 Collecting and Analyzing Database Performance Information

PostgreSQL provides several tools and views for collecting and analyzing performance information, such as:

- **pg_stata_user_tables**: This view provides statistical information about tables, including the number of rows, disk space usage, and the last autovacuum and autoanalyze times.
- **pg_stat_user_indexes**: This view provides statistical information about indexes, including the number of index scans and the index size.
- **pg_stat_database**: This view provides statistics about database-level activity, such as the number of transactions, tuples read and written, and blocks read and written.
- **EXPLAIN** and **EXPLAIN ANALYZE**: These commands provide execution plans for SQL queries, including the estimated and actual costs, and information about the operations performed.

We can use these tools to analyze query performance, identify bottlenecks, and optimize our database for better efficiency.

Here's an example of using **EXPLAIN ANALYZE** to analyze the performance of a query:

```

1 EXPLAIN ANALYZE
2 SELECT h.horseName, o.lname, o.fname
3 FROM Horse h
4 JOIN Owns ow ON h.horseId = ow.horseId
5 JOIN Owner o ON ow.ownerId = o.ownerId
6 WHERE h.age > 3;

```

This will display the execution plan for the query, including the actual time and cost for each operation, providing insights into potential performance bottlenecks. See Figure 11, for the analysis of execution plan for the Listing 10.