

# Homework 3 - Operating Systems (ICS431)

Alfaifi, Ammar - 201855360

## 1 Modified Fibonccci

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 #define MAX_SEQUENCE 100
6 #define NUM_THREADS 2
7
8 int fib_sequence[MAX_SEQUENCE];
9 int last_idx = -1; // initialize last index to -1
10
11 pthread_mutex_t mutex;
12
13 void *fibonacci(void *arg) {
14     int n = *((int *)arg);
15     fib_sequence[0] = 0;
16     fib_sequence[1] = 1;
17
18     for (int i = 2; i < n; i++) {
19         pthread_mutex_lock(&mutex);
20         fib_sequence[i] = fib_sequence[i - 1] + fib_sequence[i - 2];
21         last_idx = i; // update last index to i
22         pthread_mutex_unlock(&mutex);
23     }
24     pthread_exit(NULL);
25 }
26
27 int main(int argc, char *argv[]) {
28     pthread_t threads[NUM_THREADS];
29
30     if (argc != 2) {
31         fprintf(stderr, "Usage: %s <Number of fib elements>\n", argv[0]);
32         exit(-1);
33     }
34     int n = atoi(argv[1]);
35
36     pthread_mutex_init(&mutex, NULL);
37
38     int rc = pthread_create(&threads[0], NULL, fibonacci, (void *)&n);
39     if (rc) {
40         printf("ERROR; return code from pthread_create() is %d\n", rc);
41         exit(-1);
42     }
43
44     while (last_idx < n - 1) { // loop until all values are updated by the child
45         pthread_mutex_lock(&mutex);
46         int idx = last_idx;
47         pthread_mutex_unlock(&mutex);
48
49         if (idx >= 0) { // print new values if there are any
50             for (int i = 0; i <= idx; i++) {
51                 printf("%d ", fib_sequence[i]);
52             }
53             printf("\n");
54         }
55     }
56
57     pthread_join(threads[0], NULL);
58
59     pthread_mutex_destroy(&mutex);
60     pthread_exit(NULL);
61 }
```

## 2 Semphaore

A list of needed variables as semaphores:

1. `patient_sem` - initialized to 1, used to ensure only one patient writes to the buffer at a time
2. `doctor_sem` - initialized to 0, used to ensure doctor waits until there is new patient information in the buffer
3. `treat_sem` - initialized to 0, used to ensure the patient waits until the doctor has written the treatment details to the buffer

A pseudocode implementation can be as following:

```
semaphore patient_sem = 1;
semaphore doctor_sem = 0;
semaphore treat_sem = 0;

void patient_process() {
    // wait in waiting room until doctor is free
    wait(doctor_sem);

    // enter doctor's office and consult doctor
    wait(patient_sem);
    consultDoctor();
    signal(patient_sem);

    // wait for doctor to treat and update buffer
    signal(doctor_sem);
    wait(treat_sem);

    // note treatment and leave doctor's office
    noteTreatment();
}

void doctor_process() {
    while (true) {
        // wait for patient to arrive and update buffer
        wait(patient_sem);
        treatPatient();
        signal(treat_sem);

        // signal patient that treatment details are in buffer
        signal(doctor_sem);
    }
}
```

A real implementation in C will be as

```
1  #include <pthread.h>
2  #include <semaphore.h>
3
4  sem_t patient_sem, doctor_sem, treat_sem;
5
6  void* patient_process(void* arg) {
7      // wait in waiting room until doctor is free
8      sem_wait(&doctor_sem);
9
10     // enter doctor's office and consult doctor
11     sem_wait(&patient_sem);
12     consultDoctor();
13     sem_post(&patient_sem);
14
15     // wait for doctor to treat and update buffer
16     sem_post(&doctor_sem);
17     sem_wait(&treat_sem);
18
19     // note treatment and leave doctor's office
20     noteTreatment();
21     pthread_exit(NULL);
22 }
23
24 void* doctor_process(void* arg) {
25     while (1) {
26         // wait for patient to arrive and update buffer
27         sem_wait(&patient_sem);
28         treatPatient();
29         sem_post(&treat_sem);
30
31         // signal patient that treatment details are in buffer
32         sem_post(&doctor_sem);
33     }
34 }
35
36 int main() {
37     // initialize semaphores
38     sem_init(&patient_sem, 0, 1);
39     sem_init(&doctor_sem, 0, 0);
40     sem_init(&treat_sem, 0, 0);
41
42     // create threads for patient and doctor processes
43     pthread_t patient_thread, doctor_thread;
44     pthread_create(&patient_thread, NULL, patient_process, NULL);
45     pthread_create(&doctor_thread, NULL, doctor_process, NULL);
46
47     // wait for threads to finish
48     pthread_join(patient_thread, NULL);
49     pthread_join(doctor_thread, NULL);
50
51     // destroy semaphores
52     sem_destroy(&patient_sem);
53     sem_destroy(&doctor_sem);
54     sem_destroy(&treat_sem);
55
56     return 0;
57 }
```

### 3 Race Conditions Problem

Here are the solution of each part:

1. The variable is `available_resources`.
2. The race condition occurs in both the `decrease_count()` and `increase_count()` functions.
3. To fix the race condition, a semaphore or mutex lock can be used to ensure mutual exclusion when accessing the `available_resources` variable. One possible implementation is as follows:

```
1  #define MAX_RESOURCES 5
2  int available_resources = MAX_RESOURCES; // Critical section
3  sem_t mutex;
4
5  int decrease_count(int count) {
6      sem_wait(&mutex);
7      if (available_resources < count) {
8          sem_post(&mutex);
9          return -1;
10     } else {
11         available_resources -= count;
12         sem_post(&mutex);
13         return 0;
14     }
15 }
16
17 void increase_count(int count) {
18     sem_wait(&mutex);
19     available_resources += count;
20     sem_post(&mutex);
21 }
```

## 4 2-Thread Matrix

The implementation of two matrices of size 40x40 with random values of 0 or 1 with two separate threads is as following:

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4
5  #define ROWS 40
6  #define COLS 40
7
8  int matrixA[ROWS][COLS];
9  int matrixB[ROWS][COLS];
10 int sumMatrix[ROWS][COLS];
11 int diffMatrix[ROWS][COLS];
12
13 void *computeSum(void *arg) {
14     for (int i = 0; i < ROWS; i++) {
15         for (int j = 0; j < COLS; j++) {
16             sumMatrix[i][j] = matrixA[i][j] + matrixB[i][j];
17         }
18     }
19     pthread_exit(NULL);
20 }
21
22 void *computeDiff(void *arg) {
23     for (int i = 0; i < ROWS; i++) {
24         for (int j = 0; j < COLS; j++) {
25             diffMatrix[i][j] = matrixA[i][j] - matrixB[i][j];
26         }
27     }
28     pthread_exit(NULL);
29 }
30
31 void printMatrix(int matrix[ROWS][COLS]) {
32     for (int i = 0; i < ROWS; i++) {
33         for (int j = 0; j < COLS; j++) {
34             printf("%d ", matrix[i][j]);
35         }
36         printf("\n");
37     }
38 }
39
40 int main() {
41     pthread_t threadSum, threadDiff;
42
43     // Initialize matrices with binary inputs
44     for (int i = 0; i < ROWS; i++) {
45         for (int j = 0; j < COLS; j++) {
46             matrixA[i][j] = rand() % 2;
47             matrixB[i][j] = rand() % 2;
48         }
49     }
50
51     // Create threads for computing sum and difference
52     pthread_create(&threadSum, NULL, computeSum, NULL);
53     pthread_create(&threadDiff, NULL, computeDiff, NULL);
54
55     // Wait for threads to complete
56     pthread_join(threadSum, NULL);
57     pthread_join(threadDiff, NULL);
58
59     printf("\nSum Matrix:\n");
60     printMatrix(sumMatrix);
61
62     printf("\nDifference Matrix:\n");
63     printMatrix(diffMatrix);
64
65     return 0;
66 }
```

# Theoretical Part

## 1 Scheduling Policies

(a) Using FCFS scheduling algorithm and completing the table, we have the following:

Process	Arrival T.	Burst T.	Turnaround T.
P1	0.0	8	8
P2	0.4	4	11.6
P3	1.0	1	12

Then the average turn around time can be found as

$$\bar{T}_t = \frac{8 + 11.6 + 12}{3} = 10.5$$

Where  $\bar{T}_t$  is the average turnaround time.

(b) Doing the same but with SJF algorithm we get

Process	Arrival T.	Burst T.	Turnaround T.
P1	0.0	8	8
P3	1.0	1	8
P2	0.4	4	12.6

So we have

$$\bar{T}_t = \frac{2(8) + 12.6}{3} = 9.5$$

(c) After waiting for 1 time unit, we get the following

Process	Arrival T.	Burst T.	Turnaround T.
P3	1.0	1	1
P2	0.4	4	4.6
P1	0.0	8	13

And we get

$$\bar{T}_t = \frac{1 + 4.6 + 13}{3} = 6.2$$